

# My iTunes Dataset

## Content Page

1. Introduction
  - 1.1 Description
  - 1.2 Purpose
  - 1.3 Variables
  - 1.4 Key steps
2. Methods / Analysis
  - 2.1 Cleaning the data & the issues
  - 2.2 Some theories
  - 2.3 Modelling approach
3. Results
  - 3.1 RMSE
  - 3.2 The outliers
  - 3.3 General trend
4. Conclusion
  - 4.1 Summary
  - 4.2 Potential impact
  - 4.3 Limitations
  - 4.4 Future work

# 1. Introduction

## 1.1 Description

iTunes is where I get my music from. Using a paid family monthly subscription like that of Spotify, I am able to listen to unlimited music. On the iTunes app, I am able to retrieve very interesting pieces of information. Below is a snapshot of the dataset in its raw form.

Name	Time	Album by Artist	Artist	Plays	Year	Genre	Last Played	Release Date	Size
50 Ways to Say Goodbye	4:08	California 37	Train	19	2012	Pop	14/5/19, 7:46 AM	13/4/12	8.8 MB
25 Rasul	5:05	Brotherhood	Raihan	2	1997	Pop	6/8/19, 7:43 AM	10/12/97	11.1 MB
24K Magic	3:47	24K Magic	Bruno Mars	3	2016	Pop	4/11/17, 7:51 PM	7/10/16	7.7 MB
22	3:52	Red (Deluxe Versi...	Taylor Swift	10	2012	Country	15/10/18, 8:08 PM	12/5/09	8.2 MB
21 Guns	5:21	21st Century Brea...	Green Day	1	2009	Alternative	22/9/19, 3:25 PM	12/5/09	10.9 MB
20 ans	3:36	L'attente (Deluxe V...	Johnny Hallyday	1	2012	Rock	1/6/19, 11:01 AM	12/11/12	7.4 MB
7-й элемент (Седьмой эл...	4:05	Майский шашлы...	Vitas	2	2001	Pop	14/6/18, 8:42 PM	1/1/01	8.3 MB

Unfortunately, I don't know how to extract this as a csv. So I copy-pasted everything in to a txt file. It kinda looks like this:

1	50 Ways to Say Goodbye	4:08	California 37	Train	19	2012	Pop	14/5/19, 7:46 AM	13/4/12	8.8 MB
2	25 Rasul	5:05	Brotherhood	Raihan	2	1997	Pop	6/8/19, 7:43 AM	10/12/97	11.1 MB
3	24K Magic	3:47	24K Magic	Bruno Mars	3	2016	Pop	4/11/17, 7:51 PM	7/10/16	7.7 MB
4	22	3:52	Red (Deluxe Version)	Taylor Swift	10	2012	Country	15/10/18, 8:08 PM	12/5/09	8.2 MB
5	21 Guns	5:21	21st Century Breakdown (Deluxe Version)	Green Day	1	2009	Alternative	22/9/19, 3:25 PM	12/5/09	10.9 MB
6	20 ans	3:36	L'attente (Deluxe Version)	Johnny Hallyday	1	2012	Rock	1/6/19, 11:01 AM	12/11/12	7.4 MB
7	7-й элемент (Седьмой элемент)	4:05	Майский шашлындос. Любимые песни	Vitas	2	2001	Pop	14/6/18, 8:42 PM	1/1/01	8.3 MB
8	7 Years	3:57	Lukas Graham	Lukas Graham	3	2015	Pop	19/11/17, 5:15 PM	16/6/15	8.1 MB
9	7 Years	3:53	7 Years - Single	LittleTranscriber	2016	Pop	16/3/16	7.8 MB	1	
10	6 at Best	3:38	Out of Time - EP	First to Eleven	1	2016	Pop	27/6/19, 7:33 AM	28/4/16	7.4 MB
11	2U (feat. Justin Bieber)	3:15	2U (feat. Justin Bieber) - Single	David Guetta	1	2017	Dance	1/9/17, 6:33 PM	9/6/17	6.7 MB
12	青花瓷	3:57	On the Run	Jay Chou	6	2007	Mandopop	12/4/19, 12:24 PM	2/11/07	8.2 MB
13	随它吧 (Chinese Mandarin Version)	3:45	Let It Go the Complete Set (From "Frozen")	Jalane Hu	12	2014	Soundtrack	11/9/19, 7:11 AM	15/4/14	7.6 MB
14	醉赤壁	4:38	林俊傑2003年-2010年作品精選集	JJ Lin	2008	Mandopop	2/10/08	9.3 MB		
15	遇见 (feat. Jeremy Monteiro)	4:28	遇见 (feat. Jeremy Monteiro) - Single	Joanna Dong	2	2017	Contemporary Jazz	12/8/19, 10:43 AM	14/4/17	9 MB
16	追光者 (電視劇《夏至未至》插曲)	3:56	追光者 (電視劇《夏至未至》插曲) - Single	Yoyo Sham	2017	Mandopop	16/6/17	8 MB		
17	轉眼 (2018自傳最終章)	6:01	轉眼 (2018自傳最終章) - Single	Mayday	1	2018	Mandopop	11/4/19, 6:46 PM	31/12/18	12.2 MB
18	說好的幸福呢	4:15	周杰伦	Jay Chou	1	2008	Mandopop	8/8/19, 10:11 AM	14/10/08	8.6 MB

After that, it was easy, since every entry is separated by tabs. So I used this code to convert the text into a dataset:

```
raw_itunes <- read.delim('Arcturus/itunes.txt', header = FALSE, sep = "\t", dec = ",")
names(raw_itunes) <- c('title', 'downloaded', 'duration', 'album', 'artist', 'plays', 'year', 'genre', 'last_pl', 'size')
raw_itunes$downloaded<-NULL
```

(I removed the 'downloaded' column because it is useless..)

## 1.2 Purpose

Honestly, I've always wanted to use this dataset for something. Looking at the variables, its hard to find a tangible quantity to predict. But I wanted to know what songs are of my liking. So I ended up choosing to predict the 'Plays' variable. So, given other columns, I would like to predict how many times I've played the song. Hopefully, using this model, I am able to create a spider that finds songs that suit me. But that's a project for another time.

## 1.3 Variables

Printed below are the variables found in the dataset:

```
names(raw_itunes)
```

```
## [1] "title"      "duration"   "album"      "artist"     "plays"
## [6] "year"      "genre"     "last_played" "release_date" "size"
## [11] "skips"
```

**title** (note: LaTeX doesn't support chinese characters so I am unable to show all the titles..)

Self-explanatory, contains the title of the song. The important thing to note here is that the title sometimes contains the names of supporting artists. After much exploring, I found out that the only places where artists can be found in the title is after the 'feat.' word. Extracting these names is for another section:

```
as.vector(raw_itunes$title[1:5])
```

```
## [1] "À la plus haute branche"
## [2] "À peu près"
## [3] "The A Team"
## [4] "The a Team"
## [5] "Abe Lincoln vs Chuck Norris (feat. Nice Peter & Epiclloyd)"
```

**duration**

This represents the duration, in minutes, of the song:

```
as.vector(raw_itunes$duration[1:10])
```

```
## [1] "4:50" "3:26" "4:42" "4:18" "2:08" "2:01" "2:28" "4:24" "3:44" "4:56"
```

**album**

The name of the album that the song is from. Ended up not using it:

```
as.vector(raw_itunes$album[1:5])
```

```
## [1] "Encore un soir"
## [2] "À peu près"
## [3] "Cover Sessions, Vol. 2"
## [4] "+"
## [5] "Abe Lincoln vs Chuck Norris (feat. Nice Peter & Epiclloyd) - Single"
```

**artist**

The name of the artist(s) that created / performed the song. Sometimes contains multiple artistes, separated by either commas or the ampersand(&) sign. Same as title in this regard:

```
as.data.frame(raw_itunes %>% group_by(artist) %>%
  summarize(n=n()) %>% arrange(desc(n)))[1:4,]
```

```
##           artist    n
## 1      Boyce Avenue 118
## 2    LittleTranscriber 100
## 3      Rucka Rucka Ali  89
## 4 Epic Rap Battles of History 81
```

## plays

Our objective. This records the number of times I've listened to this song:

```
as.data.frame(raw_itunes %>% arrange(desc(plays)) %>% select(title, plays))[1:4,]
```

```
##              title plays
## 1           The Nights   106
## 2 Down (feat. Lil Wayne)   98
## 3      Rewrite the Stars   96
## 4    Alexander Hamilton   95
```

## genre

This records the genre that the song belongs in. Note that there are many different types of genres recorded, with some values missing:

```
as.data.frame(raw_itunes %>% group_by(genre) %>%
  summarize(n=n()) %>% arrange(desc(n)))[1:4,]
```

```
##           genre    n
## 1           Pop  784
## 2         Comedy  253
## 3 Singer/Songwriter 202
## 4      Alternative  167
```

## last\_played

This records the last time I've played this song. It comes as a string. If it is blank, it means I haven't played the song before:

```
as.vector(raw_itunes$last_played)[1:10]
```

```
## [1] "4/10/18, 1:10 PM" "14/1/20, 6:37 PM" "2/3/19, 10:37 AM"
## [4] "29/11/17, 4:44 AM" "4/11/19, 6:47 PM" "3/9/19, 7:13 AM"
## [7] "9/8/19, 12:36 PM" "24/6/19, 7:21 AM" "24/7/19, 6:51 PM"
## [10] "3/12/19, 8:37 PM"
```

## release\_date

This records the date that the song is released. It also comes as a string. Similar to last\_played, there are missing values:

```
as.vector(raw_itunes$release_date)[1:10]
```

```
## [1] "26/8/16" "6/10/17" "2/1/12" "7/2/10" "15/12/11" "8/2/13"
## [7] "17/9/13" "6/11/15" "17/5/18" "8/4/82"
```

## size

This records the file size of the song, aka how much space it takes up on my phone:

```
as.vector(raw_itunes$size)[1:10]
```

```
## [1] "9.7 MB" "7.1 MB" "9.3 MB" "8.7 MB" "4.7 MB" "4.4 MB" "5.3 MB"
## [8] "8.9 MB" "8.5 MB" "10.6 MB"
```

## skips

This records the number of times I've skipped the song, aka my annoyance with the song.

```
as.data.frame(raw_itunes %>% arrange(desc(skips)) %>% select(title, skips))[1:4,]
```

```
##      title skips
## 1    Hello     9
## 2 All of Me     7
## 3    Sorry     7
## 4 Good Time     6
```

## 1.4 Key Steps

- Things I need to do:
  - Convert 'last\_played' and 'release\_date' into datetime strings
  - Extract artist names from 'title' and 'artist'
  - Convert 'duration' and 'size' into their appropriate units
  - Clean data
  - Replace empty data with fake ones

## 2. Methods / Analysis

### 2.1 Cleaning the data

This data has a few issues. Firstly, there are quite a few missing values.

#### Missing values in skips, plays, genre

For the numeric ones like 'skips' and 'plays', missing values are due to zeroes. For 'genre', after some searching, I've found that missing values are all 'Instrumental' in nature. So using this code, I am able to replace them:

```
raw_itunes$plays[is.na(raw_itunes$plays)] <- 0
raw_itunes$skips[is.na(raw_itunes$skips)] <- 0
raw_itunes$genre[raw_itunes$genre==''] <- 'Instrumental'
```

#### Reassigning genre

There are an unnecessarily large number of genres in this dataset, including the low-count ones like:

```
as.data.frame(raw_itunes %>% group_by(genre) %>%
  summarize(n=n()) %>% arrange(n))[1:4,]
```

```
##           genre n
## 1 Adult Contemporary 1
## 2           Bollywood 1
## 3           Brazilian 1
## 4           Chinese Rock 1
```

As such, I needed a way to reassign the genres to encompass a smaller variety. I used this code to map the old to the new genres:

```
unique_genres <- unique(raw_itunes$genre)
genres_that_i_want <- c('Non-Asian', 'Asian', 'Pop', 'Jazz', 'Dance-ish', 'Rock-ish', 'For Kids',
  'Soundtrack', 'Rap', 'Covers', 'Instrumental', 'Alternative', 'Comedy', 'Misc')
assigned_genres <- genres_that_i_want[c(1,10,13,5,12,11,6,1,7,3,
  5,9,8,11,3,12,3,12,2,14,
  11,5,6,11,14,1,4,14,9,3,
  9,6,3,2,1,1,5,11,14,6,
  4,9,2,4,2,11,2,2,1,8,
  3,1,2,4)]

genre_map <- setNames(assigned_genres, unique_genres)
raw_itunes$genre <- genre_map[as.vector(raw_itunes$genre)]
```

So now, the least common genres have the following numbers of songs:

```
as.data.frame(raw_itunes %>% group_by(genre) %>%
  summarize(n=n()) %>% arrange(n))[1:4,]
```

```
##           genre n
## 1           Jazz 4
## 2           Asian 10
## 3 Soundtrack 33
## 4   For Kids 41
```

## Reformatting duration

The ‘duration’ column needs to be reformatted. It is currently of the format ‘MM:SS’, where the song is MM minutes and SS seconds long. So, I created a new column, labelled ‘time\_in\_seconds’ which records the duration in terms of seconds:

```
raw_itunes$time_in_seconds <- sapply(as.vector(raw_itunes$duration), function(x){
  as.numeric(strsplit(x,':')[[1]][1])*60+as.numeric(strsplit(x,':')[[1]][2])
})
# Remove invalid / zero length durations
raw_itunes$time_in_seconds[is.na(raw_itunes$time_in_seconds)] <- 0
raw_itunes <- subset(raw_itunes, time_in_seconds!=0)
```

So now, we can see the longest songs:

```
as.data.frame(raw_itunes %>% arrange(desc(time_in_seconds)) %>% select(title, time_in_seconds))[1:4,]

##                                                                 title
## 1 Four_Chord_Mega-Medley_Alان_Walker_Imagine_Dragons_Avril_Lavigne_The_Script__more
## 2                                                                 H20901
## 3                                                                 Albuquerque
## 4          Trapped In the Drive-Thru (Parody of Trapped In the Closet By R. Kelly)
##   time_in_seconds
## 1             1627
## 2             1560
## 3              683
## 4             651
```

## Reformatting size

The ‘size’ column also needs to be reformatted. Currently, it is either of the form ‘x KB’ or ‘x MB’. I want to convert everything into a standard format. So I used the following code to convert everything to KB as a new column:

```
library(tidyr)
raw_itunes$kb <- sapply(as.vector(raw_itunes$size),function(x){
  if (grepl('MB',x)) {
    as.numeric(gsub(' MB', '',x))*1000
  } else {
    as.numeric(gsub(' KB', '',x))
  }
})
```

## Reformatting dates

‘release\_date’ and ‘last\_played’ are still in string format, which is not very useful. So, I used this code to convert them into datetime formats:

```
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##      date
raw_itunes$clean_release_date <- as.Date(raw_itunes$release_date, format='%d/%m/%y')
raw_itunes$clean_last_played <-as.Date(raw_itunes$last_played, format='%d/%m/%y, %I:%M %p')
```

## Faking the data

I found that the two dates columns mentioned have quite a lot of missing data. At first I thought that this wasn't an issue. But later on I found that this causes a lot of problems as I will have very few predicting variables. As such, I decided to create fake data. I first retrieve the means and sd of each column (non-NAs) and then I use `rtruncnorm` to generate new dates that follow the distribution. Somehow, it manages to work. So for example, the average and sd of 'clean\_release\_date' are shown below:

```
print(mean(raw_itunes$clean_release_date[!is.na(raw_itunes$clean_release_date)]))

## [1] "2013-01-28"

print(sd(raw_itunes$clean_release_date[!is.na(raw_itunes$clean_release_date)]))

## [1] 2469.313
```

## Further cleaning of release date

I also found out another problem. For the release date, the format was `%d/%m/%y`, which meant that the year was reported as double digits. As such, some songs early in the 20th century were wrongly reported as being in the 21st century:

```
raw_itunes %>% filter(clean_release_date > Sys.Date()) %>% select(title, clean_release_date)

##           title clean_release_date
## 1 The Elements      2053-01-01
## 2 La Cucaracha      2062-01-01
```

So that is fixed with this code:

```
year(raw_itunes$clean_release_date[raw_itunes$clean_release_date>current_date]) <-
+year(raw_itunes$clean_release_date[raw_itunes$clean_release_date>current_date])-100
```

## Extracting all artists

The 'artists' column is far from complete. Multiple artists in the same song are separated by commas. Same goes for the title, where the word 'feat.' appears in many songs. As such, I decided that I want to create a separate row for each artist in the song. Firstly, I created the column 'songId' in order to not lose track of the song in question:

```
raw_itunes$songId<-c(1:length(raw_itunes$title))
```

Next is just a sequence of `str_split`, `str_replace_all`, `gsub`, and `sapply` to get what I want (too long to show). So now my top artist list looks like this:

```
as.data.frame(clean_itunes %>% group_by(artist) %>%
  summarize(n=n()) %>% arrange(desc(n)))[1:4,]

##           artist    n
## 1      Boyce Avenue 120
## 2  LittleTranscriber 100
## 3      Rucka Rucka Ali  89
## 4 Epic Rap Battles of History 81
```

(Not much difference, I know, but trust me on this. BTW since this is a major change, I renamed the dataset 'clean\_itunes')



## 2.2 Some theories

I have some theories on what variables affect the number of plays in a song.

### info\_density

First one is info\_density, which is how dense the song is. Maybe the denser the music, the more I like it?? So I used the 'kb' and 'time\_in\_seconds' to create this column.


```
clean_itunes <- clean_itunes%>%mutate(info_density = kb/time_in_seconds)
```

And now, let's look at the correlation coefficient and the distribution of the info\_density:

```
library(ggplot2)
print(cor(clean_itunes$plays,clean_itunes$info_density))
```

```
## [1] -0.01827507
```

```
clean_itunes %>% group_by(genre) %>% filter(info_density<100) %>%
ggplot(aes(x=info_density,y=plays,color=genre))+geom_point()
```



Doesn't quite work, but I'll still include it.

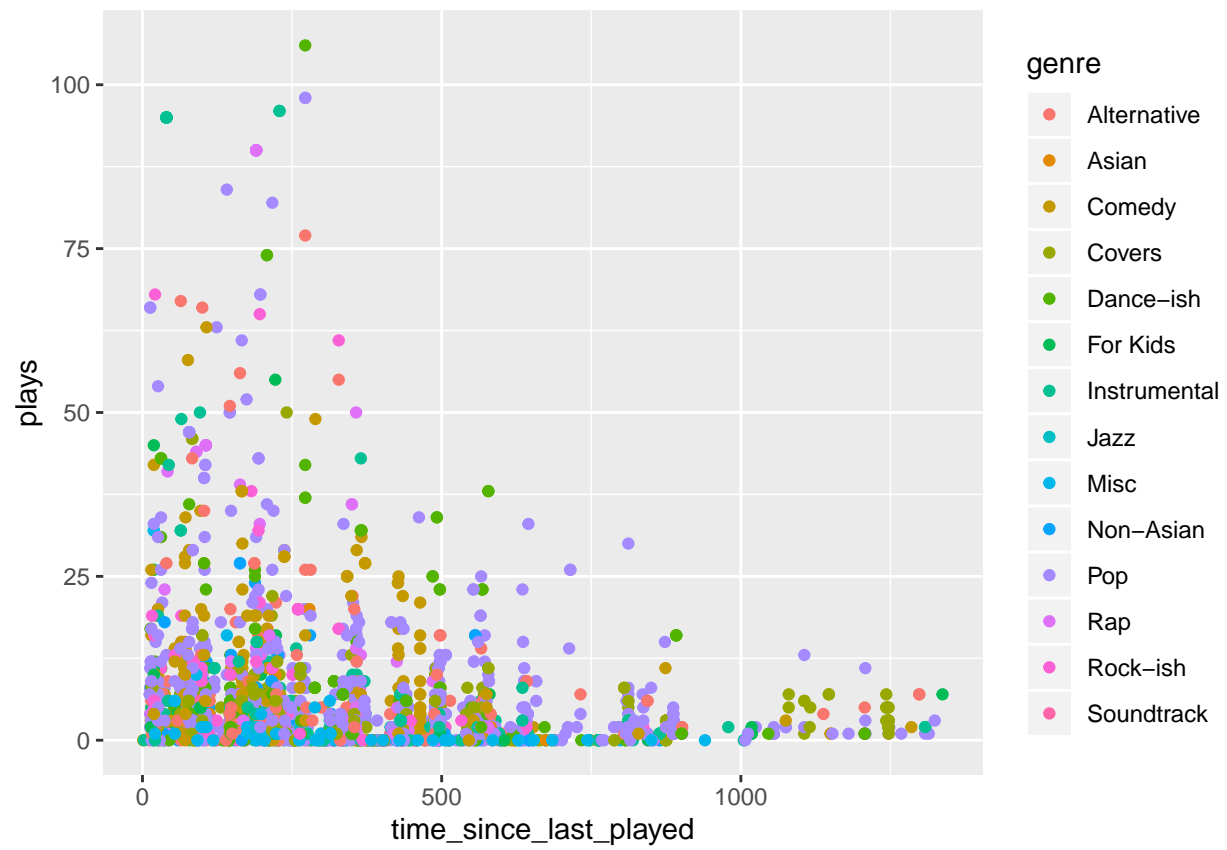
### time\_since\_last\_played

If it's been a long time since I played a song, it usually means that i rarely play it:

```
clean_itunes$time_since_last_played <-
  as.numeric(current_date) - as.numeric(clean_itunes$clean_last_played)
print(cor(clean_itunes$time_since_last_played,clean_itunes$plays))
```

```
## [1] -0.2007519
```

```
clean_itunes %>% group_by(genre) %>%
  ggplot(aes(x=time_since_last_played,y=plays,color=genre))+geom_point()
```




```
### song_age_when_last_played
```

As the name suggests:

```
clean_itunes$song_age <- as.numeric(current_date) - as.numeric(clean_itunes$clean_release_date)
clean_itunes$song_age_when_last_played <- clean_itunes$song_age - clean_itunes$time_since_last_played
print(cor(clean_itunes$song_age_when_last_played, clean_itunes$plays))
```

```
## [1] -0.06286705
```

```
clean_itunes %>% group_by(genre) %>%
ggplot(aes(x=song_age_when_last_played, y=plays, color=genre))+geom_point()
```



### Variables bound to artist

These are variables I can only define within the train set because I'm not supposed to use test data as predictors to avoid overfitting. They are created by first grouping by artist and then do some other operations:

#### artist\_time\_spent

Finds out how much time in total I've spent listening to the artist. `sum(plays*time_in_seconds)`

#### artist\_plays

How many times I've played one of their songs. `sum(plays)`

#### artist\_skips

How many times I've skipped one of their songs. `sum(skips)`

**artist\_total\_songs**

How many of their songs I have. `n()`

**artist\_avg\_time\_spent**

Average time spent on their songs. `mean(plays*time_in_seconds)`

**artist\_avg\_plays**

Average number of times I've played their songs. `mean(plays)`

**artist\_avg\_skips**

Average number of times I've skipped their songs. `mean(skips)`

**Variables bound to genre**

Exactly the same as that for artist.

## 2.3 Modelling approach

After doing the cleaning, I separated the data into an initial 0.9,0.1 split using `createDataPartition`. A `train_set` and a validation set. I used a `sapply` loop to loop through a vector 1:100. Within that loop, from the `train_set`, I separated that further into a 0.9,0.1 split. An `itunes` set and a `test_set`. I performed the artist-wise and the genre-wise operations on the `itunes` set. Similar to the `cor()` functions I've done earlier, I then found out what variables are most correlated to the number of plays, choosing the best 10 as predictors:

```
df_of_cors <- as.data.frame(cor(itunes[sapply(itunes,is.numeric)],itunes$plays))
names(df_of_cors) <- c('R')
df_of_cors$variable <- row.names(df_of_cors)
variables_to_train <- df_of_cors[order(-abs(df_of_cors$R)),][2:11,]$variable
```

I then used three different training models from the `caret` package:

```
TrainData <- itunes %>% select(variables_to_train)
TrainClasses <- itunes$plays
model_1 <- train(TrainData,TrainClasses,method='glm',family='gaussian',
  tuneLength=10,trControl=trainControl(method='cv'))
model_2 <- train(TrainData,TrainClasses,method='glm',family='quasipoisson',
  tuneLength=10,trControl=trainControl(method='cv'))
model_3 <- train(TrainData,TrainClasses,method='gamLoess',
  tuneLength=10,trControl=trainControl(method='cv'))
```

After that, I merged the artist-wise and genre-wise data from the `itunes` dataset into the `test_set`. So now, the `test_set` has the necessary predictors. After some cleaning, I predicted the `test_set` plays. I picked the best model out of all of them (best RMSE) and if the RMSE is less than 10, I would use that model to predict the validation plays. And out of all these validation predictions, I averaged them to get my final prediction. This process took like 2 hours to complete so I'm not going to show it here.

## 3. Results

### 3.1 RMSE

I experimented with many regression models. Some had issues, other didn't. Some churned out abysmal RMSEs, others gave adequate ones. Out of the 100 times I've split the data, only 18 of them reported RMSEs less than 10. Quite sad, but I'll have to deal with it. In the end, the final RMSE I got was quite bad, at 12.06523.

## 3.2 The outliers

I decided to find out what were causing the values to be way off. Using this code, I managed to find out which rows gave the lowest errors:

```
> as.data.frame(validation %>% arrange(error))[1:10,]
  title duration album artist plays year genre
1 Tunak Tunak Tun 5:03 Tunak Tunak Tun Daler Mehndi 2 1998 Asian
2 Hotline Bling 2:58 Hotline Bling - Single LittleTunscraper 2 2015 Pop
3 Lonely Together (feat. Rita Ora) 3:02 Avici (01) - EP Avicii 7 2017 Dance-ish
4 Don't You Worry Child 4:06 Wonders The Piano Guys 4 2014 Instrumental
5 Shallow 3:36 A Star Is Born Soundtrack Lady Gaga 3 2018 Instrumental
6 You Are the Reason 3:24 Only Human (Deluxe) Calum Scott 3 2017 Pop
7 This Is What It Feels Like (feat. Trevor Guthrie) 3:23 Intense (Bonus Track Version) Trevor Guthrie 1 2013 Dance-ish
8 Jumper 1:43 Community Favorites Waterflame 4 2015 Dance-ish
9 Collide (Original) 4:09 Stop All the World Now Howie Day 4 2003 Pop
10 Wavin' Flag 3:41 Troubadour (Champion Edition) K'naan 5 2009 Rap

  last_played release_date size skips time_in_seconds kb info_density clean_release_date clean_last_played time_since_last_played
1 17/7/18, 5:52 PM 30/9/98 10.4 MB 0 303 10400 34.32343 1998-09-30 2018-07-17 579
2 15/9/17, 2:19 AM 30/11/15 6.1 MB 0 178 6100 34.26966 2015-11-30 2017-09-15 884
3 25/7/19, 6:57 AM 10/8/17 6.5 MB 0 182 6500 35.71429 2017-08-10 2019-07-25 206
4 24/9/19, 7:08 AM 6/10/14 8.4 MB 0 246 8400 34.14634 2014-10-06 2019-09-24 145
5 12/3/19, 7:00 AM 5/10/18 7.4 MB 0 216 7400 34.25926 2018-10-05 2019-03-12 341
6 9/10/18, 10:54 AM 17/11/17 6.9 MB 0 204 6900 33.82353 2017-11-17 2018-10-09 495
7 30/10/16, 9:07 AM 8/4/13 7.2 MB 0 203 7200 35.46798 2013-04-08 2016-10-30 1204
8 2/3/19, 9:53 AM 14/2/15 3.8 MB 0 103 3800 36.89320 2015-02-14 2019-03-02 351
9 28/11/19, 9:23 PM 7/10/03 8.7 MB 0 249 8700 34.93976 2003-10-07 2019-11-28 80
10 14/11/19, 7:25 AM 24/2/09 7.5 MB 0 221 7500 33.93665 2009-02-24 2019-11-14 94

  song_age song_age_when_last_played songId prediction error
1 7809 7230 1934 1.982214 0.01778602
2 1539 655 776 1.981429 0.01857054
3 920 714 1089 6.975913 0.02408744
4 1959 1814 473 4.029127 0.02912650
5 499 158 1605 3.098127 0.09812660
6 821 326 2141 3.146440 0.14643966
7 2505 1301 1849 1.152513 0.15251295
8 1828 1477 961 3.833532 0.16646756
9 5976 5896 343 4.174189 0.17418885
10 4009 3915 1995 5.175586 0.17558569
```

Out of this, I noticed one thing: All the play values are small

Which kinda makes sense I guess, as it is easier to guess a small number more accurately than a large number. Also, I found out which rows gave the highest errors:

```
> as.data.frame(validation %>% arrange(desc(error)))[1:10,]
  title duration album
1 The Way I Are (Dance With Somebody) [feat. Lil Wayne] 3:08 The Way I Are (Dance With Somebody) [feat. Lil Wayne] - Single
2 Last Hurrah 2:30 Last Hurrah - Single
3 Riff Off 4:49 Pitch Perfect 3 (Original Motion Picture Soundtrack)
4 Tour the World 8:34 Brain Beats 2
5 Scared to Be Lonely 3:41 Scared to Be Lonely - Single
6 Just a Dream 3:58 5.0
7 On My Way 3:14 On My Way - Single
8 Love Story 3:55 Fearless (Platinum Edition)
9 Journey Back to You (feat. NerdOut) 3:56 Young Unprofessionals (Acoustic)
10 Journey Back to You (feat. NerdOut) 3:56 Young Unprofessionals (Acoustic)

  artist plays year genre last_played release_date size skips time_in_seconds kb info_density clean_release_date
1 Lil Wayne 3 2017 Pop 4/2/20, 5:26 PM 19/5/17 6.7 MB 0 188 6700 35.63830 2017-05-19
2 Bebe Rexha 82 2019 Pop 17/7/19, 6:44 PM 15/2/19 5.3 MB 0 150 5300 35.33333 2019-02-15
3 Evermoist 47 2017 Pop 3/12/19, 6:10 PM 15/12/17 9.9 MB 1 289 9900 34.25606 2017-12-15
4 Renald Francoeur 55 2013 For Kids 12/7/19, 6:57 AM 2/7/13 16.9 MB 3 514 16900 32.87938 2013-07-02
5 Martin Garrix 37 2017 Dance-ish 23/5/19, 1:39 PM 27/1/17 7.6 MB 0 221 7600 34.38914 2017-01-27
6 Nelly 39 2010 Rap 9/9/19, 6:58 AM 17/8/10 8.1 MB 1 238 8100 34.03361 2010-08-17
7 Alan Walker 43 2019 Dance-ish 19/1/20, 9:48 PM 21/3/19 6.7 MB 2 194 6700 34.53608 2019-03-21
8 Taylor Swift 30 2009 Pop 29/11/17, 11:44 AM 8 MB 2 235 8000 34.04255 2013-01-05
9 Ben Schuller 29 2019 Covers 27/6/19, 9:39 AM 26/7/19 7.9 MB 0 236 7900 33.47458 2019-07-26
10 NerdOut 29 2019 Covers 27/6/19, 9:39 AM 26/7/19 7.9 MB 0 236 7900 33.47458 2019-07-26

  clean_last_played time_since_last_played song_age song_age_when_last_played songId prediction error
1 2020-02-04 12 1003.000 991.000 1996 139.370310 136.37031
2 2019-07-17 214 366.000 152.000 1007 4.742559 77.25744
3 2019-12-03 75 793.000 718.000 1499 6.403177 40.59682
4 2019-07-12 219 2420.000 2201.000 1907 17.018560 37.98144
5 2019-05-23 269 1115.000 846.000 1567 4.103928 32.89607
6 2019-09-09 160 3470.000 3310.000 963 6.169941 32.83006
7 2020-01-19 28 332.000 304.000 1318 13.862669 29.13733
8 2017-11-29 809 2597.333 1788.333 1121 4.730965 25.26904
9 2019-06-27 234 205.000 -29.000 959 3.821549 25.17845
10 2019-06-27 234 205.000 -29.000 959 5.025656 23.97434
```

Out of this, I noticed two things: 1. (Almost) All the play values are large 2. I rarely listen to these artists So I guess it kinda makes sense that the model would underestimate most of these songs.

And then I thought to myself: “Only the sith deals in absolutes” so I turned this around and decided to measure errors percentage-wise. I used the following metric as my new error column:

```
validation <- validation %>% mutate(error=abs(plays/prediction-1))
```

And I got these results for the best rows:


```
> as.data.frame(validation %>% arrange(error))[1:10,]
  title duration album artist plays year genre last_played release_date size skips
1 Lonely Together (feat. Rita Ora) 3:02 Avicii (01) - EP Avicii 7 2017 Dance-ish 25/7/19, 6:57 AM 10/8/17 6.5 MB 0
2 Don't You Worry Child 4:06 Wonders The Piano Guys 4 2014 Instrumental 24/9/19, 7:08 AM 6/10/14 8.4 MB 0
3 Tunak Tunak Tun 5:03 Tunak Tunak Tun Daler Mehndi 2 1998 Asian 17/7/18, 5:52 PM 30/9/98 10.4 MB 0
4 Hotline Bling 2:58 Hotline Bling - Single LittleTranscriber 2 2015 Pop 15/9/17, 2:19 AM 30/11/15 6.1 MB 0
5 Shallow 3:36 A Star Is Born Soundtrack Lady Gaga 3 2018 Instrumental 12/3/19, 7:00 AM 5/10/18 7.4 MB 0
6 Wavin' Flag 3:41 Troubadour (Champion Edition) K'naan 5 2009 Rap 14/11/19, 7:25 AM 24/2/09 7.5 MB 0
7 Collide (Original) 4:09 Stop All the World Now Howie Day 4 2003 Pop 28/11/19, 9:23 PM 7/10/03 8.7 MB 0
8 Jumper 1:43 Community Favorites Waterflame 4 2015 Dance-ish 2/3/19, 9:53 AM 14/2/15 3.8 MB 0
9 You Are the Reason 3:24 Only Human (Deluxe) Calum Scott 3 2017 Pop 9/10/18, 10:54 AM 17/11/17 6.9 MB 0
10 My Happy Ending 4:02 Under My Skin Avril Lavigne 8 2004 Pop 21/5/19, 6:07 PM 25/5/04 8.6 MB 1
time_in_seconds kb info_density clean_release_date clean_last_played time_since_last_played song_age song_age_when_last_played songId prediction error
1 182 6500 35.71429 2017-08-10 2019-07-25 206 920 714 1089 6.975913 0.003452944
2 246 8400 34.14634 2014-10-06 2019-09-24 145 1959 1814 473 4.029127 0.007228987
3 303 10400 34.32343 1998-09-30 2018-07-17 579 7809 7230 1934 1.982214 0.008972806
4 178 6100 34.26966 2015-11-30 2017-09-15 884 1539 655 776 1.981429 0.009372295
5 216 7400 34.25926 2018-10-05 2019-03-12 341 499 158 1605 3.098127 0.031672884
6 221 7500 33.93665 2009-02-24 2019-11-14 94 4009 3915 1995 5.175586 0.033925763
7 249 8700 34.93976 2003-10-07 2019-11-28 80 5976 5896 343 4.174189 0.041729988
8 103 3800 36.89320 2015-02-14 2019-03-02 351 1828 1477 961 3.833532 0.043424064
9 204 6900 33.82353 2017-11-17 2018-10-09 495 821 326 2141 3.146440 0.046541385
10 242 8600 35.53719 2004-05-25 2019-05-21 271 5745 5474 1227 8.411011 0.048865846
```

And these for the worst rows:

```
> as.data.frame(validation %>% arrange(desc(error)))[1:10,]
  title duration album artist plays year genre last_played release_date size skips
1 Last Hurrah 2:30 Last Hurrah - Single Bebe Rexha 82 2019
2 Scared to Be Lonely 3:41 Scared to Be Lonely - Single Martin Garrix 37 2017
3 Lean On (feat. MØ) 2:57 Peace Is the Mission DJ Snake 16 2015
4 Journey Back to You (feat. NerdOut) 3:56 Young Unprofessionals (Acoustic) Ben Schuller 29 2019
5 Riff Off 4:49 Pitch Perfect 3 (Original Motion Picture Soundtrack) Evermoist 47 2017
6 The Hardest Karaoke Song in the World (feat. Steindi Jr.) 3:20 The Hardest Karaoke Song in the World (feat. Steindi Jr.) - Single Inspired by Iceland 27 2017
7 Love Story 3:55 Fearless (Platinum Edition) Taylor Swift 30 2009
8 Just a Dream 3:58 5.0 Nelly 39 2010
9 Down (feat. DJ Not Nice & Lil Wang) 4:02 Rucka's World DJ Not Nice 27 2012
10 Journey Back to You (feat. NerdOut) 3:56 Young Unprofessionals (Acoustic) NerdOut 29 2019
genre last_played release_date size skips time_in_seconds kb info_density clean_release_date clean_last_played time_since_last_played song_age
1 Pop 17/7/19, 6:44 PM 15/2/19 5.3 MB 0 150 5300 35.33333 2019-02-15 2019-07-17 214 366.000
2 Dance-ish 23/5/19, 1:39 PM 27/1/17 7.6 MB 0 221 7600 34.38914 2017-01-27 2019-05-23 269 1115.000
3 Dance-ish 10/9/17, 12:42 PM 2/3/15 6.3 MB 0 177 6300 35.59322 2015-03-02 2017-09-10 889 1812.000
4 Covers 27/6/19, 9:39 AM 26/7/19 7.9 MB 0 236 7900 33.47458 2019-07-26 2019-06-27 234 205.000
5 Pop 3/12/19, 6:10 PM 15/12/17 9.9 MB 1 289 9900 34.25606 2017-12-15 2019-12-03 75 793.000
6 Non-Asian 9/9/19, 7:21 PM 14/10/17 6.9 MB 0 200 6900 34.50000 2017-10-14 2019-09-09 160 855.000
7 Pop 29/11/17, 11:44 AM 8 MB 2 235 8000 34.04255 2013-01-05 2017-11-29 809 2597.333
8 Rap 9/9/19, 6:58 AM 17/8/10 8.1 MB 1 238 8100 34.03361 2010-08-17 2019-09-09 160 3470.000
9 Comedy 12/2/19, 8:05 PM 11/9/12 8.8 MB 0 242 8800 36.36364 2012-09-11 2019-02-12 369 2714.000
10 Covers 27/6/19, 9:39 AM 26/7/19 7.9 MB 0 236 7900 33.47458 2019-07-26 2019-06-27 234 205.000
song_age_when_last_played songId prediction error
1 152.000 1007 4.742559 16.290243
2 846.000 1567 4.103928 8.015753
3 923.000 1022 2.108041 6.589984
4 -29.000 959 3.821549 6.588546
5 718.000 1499 6.403177 6.340106
6 695.000 712 4.070530 5.633043
7 1788.333 1121 4.730965 5.341202
8 3310.000 963 6.169941 5.320968
9 2345.000 483 4.511004 4.985364
10 -29.000 959 5.025656 4.770391
```


### 3.3 General trend

Here is the distribution of the errors:



,which shows that the majority of the play values were estimated quite well, especially for the miscellaneous genres. However, the Non-Asian songs showed much greater errors, probably because I don't listen to it often. Here is another plot:





,which shows the general trend that the more I listen to a song, the higher the error is for the song. This is also shown by the correlation coefficient of 0.7911089 between the error and the number of plays.

## 4. Conclusion

### 4.1 Summary

I wanted to find out if there was a way to predict the number of times I listen to a song based on all the other variables in my itunes dataset. I found out that while it was possible, the predictions are way off. This effect is severely compounded for one-hit wonder songs, aka when I only have one song from a particular artist. Other than that, in hindsight, many of the variables that I have been using are not optimal. What I intended at the start was to create a spider program that will use my model to pick out potential songs that I would listen to. However, when I use the 'skips' and 'last\_played' columns, the big assumption would be that I have already listened to the song, beating the whole purpose of the model. ## 4.2 Potential Impact My original intention flawed, I can't really think of any other use for this model, other than to compare it with my friends'. Kinda shameful that songs appearing in the dataset are 'Tunak Tunak Tun' and 'Big and Chunky'. ## 4.3 Limitations Already mentioned ## 4.4 Future work I'll try to find a more universal dataset that will improve my prediction model. Reduce my reliance on historical data ('skips', 'last\_played', etc..). And once I re-run this model on this newfound dataset, I will create my spider program. Apart from this I guess I can just compile as many friends' datasets as I can to find out what songs of theirs I would like, but seeing as though people rarely use Apple Music, this is not a viable option.